Parallelization of MDAnalysis

Egor Marin, Google Summer of Code contributor

marinegor@fastmail.com

Motivation

- `T(insight) = T(simulation) + T(analysis)`
- `pmda` : 30 ⭐, 21 🍴 (forks)
- all MDAnalysis instances are parallelizable (courtesy @yuxuanzhuang)

Goals

Change `AnalysisBase` in a way that...

1. Introduces multiple job support
   - local -- using multiprocessing?
   - non-local -- using dask?
2. Does not break existing code (subclasses, dependencies, etc)
3. Allows customization
   - add your own execution backend
   - write your own parallelizable classes

3

Sneak peak: results

```
>>> from MDAnalysis.analysis.rms import RMSD
>>> u, ref = ...

>>> % time long_run = RMSD(u, ref)
Wall time: 4 min

>>> % time long_run = RMSD(u, ref, backend='multiprocessing', n_workers=8)
Wall time: 50 sec
```

Sneak peak: results

```
>>> from MDAnalysis.analysis.rdf import InterRDF
>>> s1, s2 = ...

>>> % time long_run = InterRDF(s1, s2)
Wall time: 1 h

>>> % time long_run = InterRDF(s1, s2, backend='multiprocessing', n_workers=8)
Wall time: 8 min
```

Implementation

General `AnalysisBase.run()` protocol:

```python
def run(self,
        start=None,
        stop=None,
        step=None,
        frames=None, ...):

    # prepare frames -- check boundaries, setup reader
    self._setup_frames(self._trajectory, start=start, stop=stop, step=step, frames=frames)

    # initialize attributes with intermediate results
    self._prepare()

    # go through trajectory
    for i, ts in enumerate(self._sliced_trajectory, ...):
        self._single_frame()

    # convert intermediate results into final ones
    self._conclude()
```

6

Implementation

split-apply-combine:

- split all frames into groups
- process each group independently
- combine results

Implementation

Additional methods of `AnalysisBase` :

- "split": `_setup_computation_groups()`
- "apply": `_compute`
- "combine": `_get_aggregator` & `ResultsGroup.merge()`
- housekeeping: `_configure_backend`

Implementation

Additional classes: `ResultsGroup`

```python
class ResultsGroup:
    def __init__(self, lookup: dict[str, Callable]): ...
    def merge(self, objects: Sequence[Results],
              require_all_aggregators: bool = True) -> Results:
        ...
```

Implementation

Additional classes: ResultsGroup

```
>> from MDAnalysis.analysis.parallel import ResultsGroup, Results
>> group = ResultsGroup(lookup={'mass': lambda floats: np.mean(floats)}
>> obj1 = Results(mass=1)
>> obj2 = Results(mass=3)
>> group.merge([obj1, obj2])
{'mass': 2.0}
```

Implementation

Additional classes: `ResultsGroup`

```python
# skip attributes you don't want to aggregate
>>> lookup = {'mass': ResultsGroup.float_mean, 'trajectory': Trajectory}
>>> group = ResultsGroup(lookup)
>>> objects = [Results(mass=1, trajectory=None), Results(mass=3, trajectory=Trajectory(...))]
>>> group.merge(objects, require_all_aggregators=False)
{'mass': 2.0}
```

Implementation

Additional classes: BackendBase

```python
from MDAnalysis.parallel import BackendBase

class CustomBackend(BackendBase):
    def __init__(self, some_resource):
        self.some_resource = some_resource

    def apply(self,
              func: Callable[T, R],
              computations: list[T]) -> list[R]:
        results = [
          self.some_resource.do_compute(func, task)
          for task in computations]
        return results
```

Implementation

Additional classes: BackendBase

```
>>> backend_instance = CustomBackend(some_resource=...)
>>> my_run = RMSD(u, ref)
>>> my_run.run(backend=backend_instance)
```

Implementation

Conditions for the `backend=...`

- `Subclass.is_paralllizable() == True`
- `backend=` `str|BackendBase`
  - if `str`, one of built-in backends
  (`multiprocessing` / `dask`)
  - if `BackendBase` subclass, explicitly say
  `unsafe=True`

14

Add to your subclass

Example: RMSD

```python
from MDAnalysis.analysis.parallel import ResultsGroup

class RMSD(AnalysisBase):
    @classmethod
    @property
    def available_backends(cls):
        return ('serial', 'multiprocessing', 'dask',)

    @classmethod
    @property
    def is_parallelizable(self):
      return True

    def _get_aggregator(self):
      return ResultsGroup(lookup={'rmsd': ResultsGroup.ndarray_vstack})
```

How fast?

18
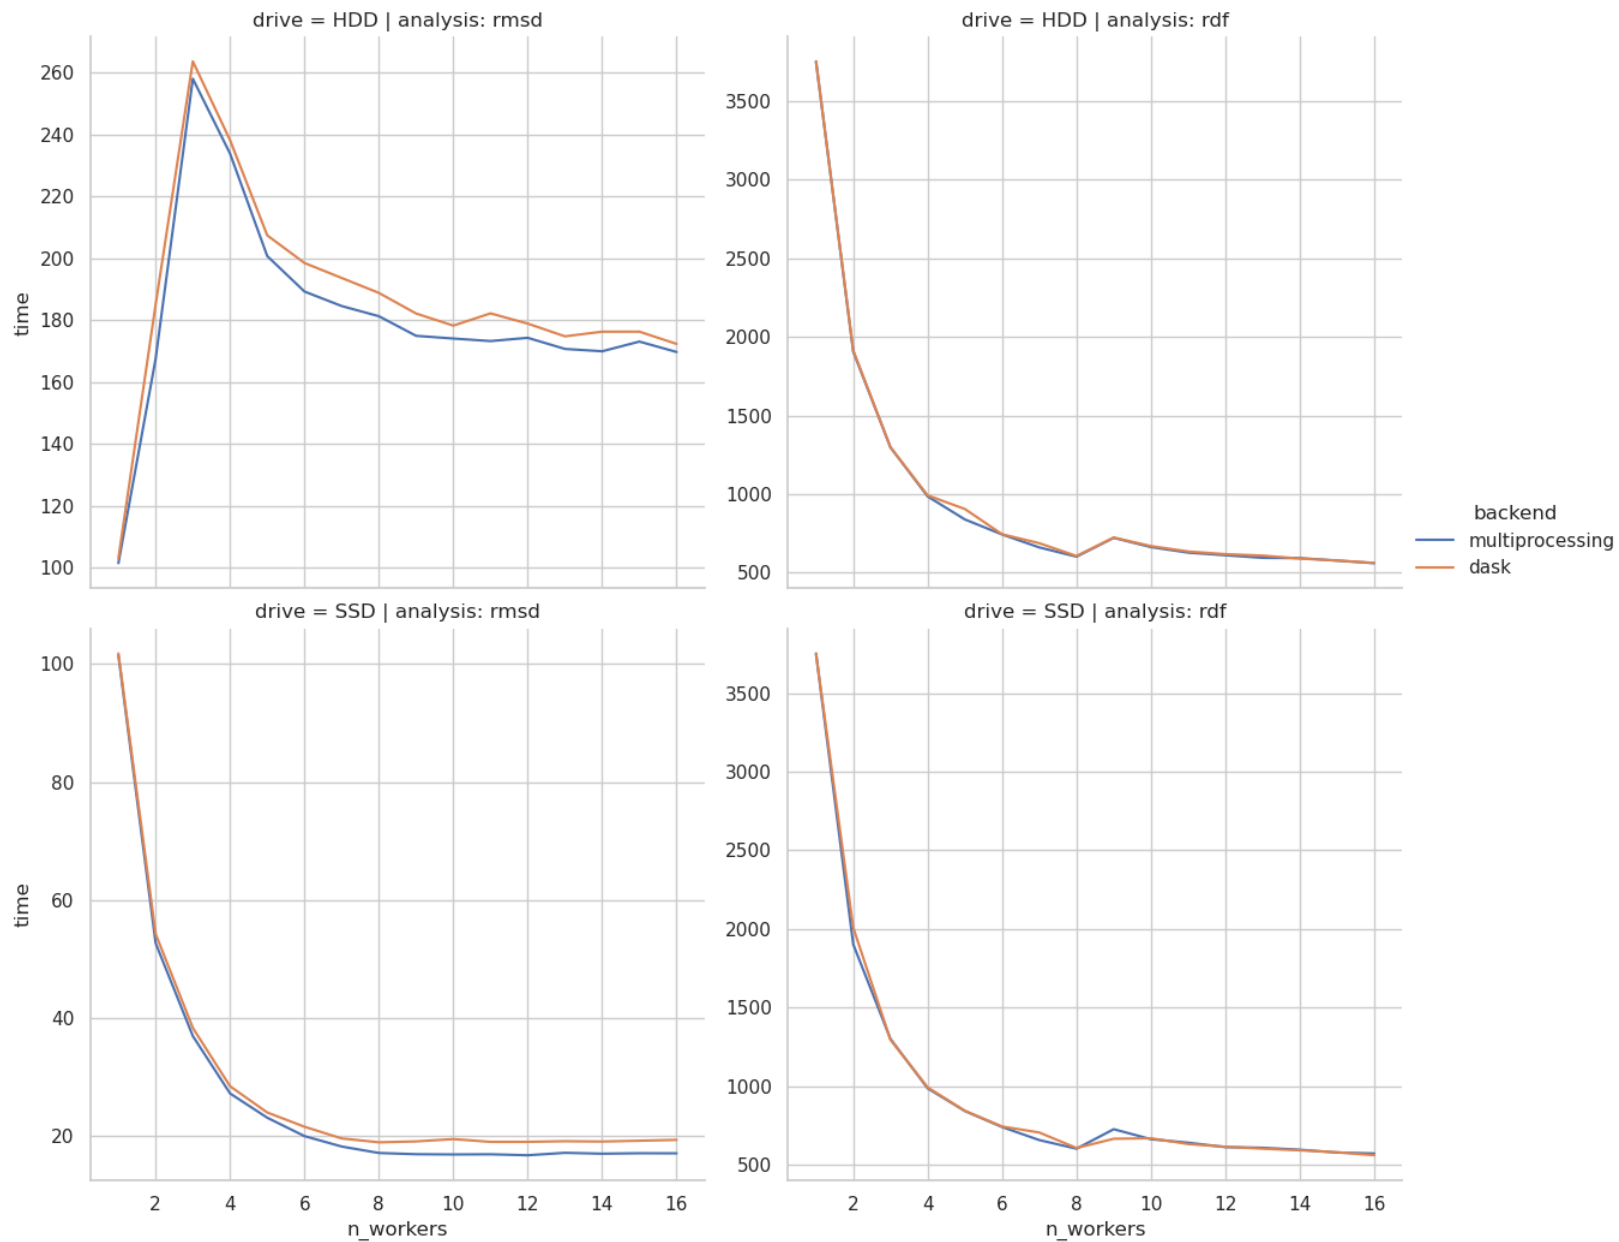
When should I use it?

When should I use it?

- `YourClass.is_parallelizable()` for your class
- `_single_frame()` slower than reading from disk

Future features

- ➡️ **SOON** add to all subclasses
  - ○ have separate issues for complicated ones?
- ➡️ **SOON** add tutorials
  - ○ how to use parallel backends
  - ○ how to implement your own backend subclasses

Future features

- 🤔 add cluster support
  - rebrand `pmda` into MDAKit
  - focus on parallel filesystems, etc
- 🤔 add shareable memory trajectories
  - pre-load trajectory & analyze it
  - efficient for practically all subclasses
- 🤔 add `AnalysisCollection` (ht @PicoCentauri)
  - one frame read exactly once

Acknowledgements

- GSOC programm
- Mentors: @yuxuanzhuang, @orbeckst, @RMeli
- all contributors🧡